

# Package ‘PFoptim’

June 28, 2022

**Type** Package

**Title** Global Stochastic Optimization using a Particle Filter Algorithm

**Version** 1.0

**Date** 2021-07-27

**Author** Mathieu Gerber

**Maintainer** Mathieu Gerber <mathieu.gerber@bristol.ac.uk>

**Description** This package implements the G-PFSO (Global Particle Filter Stochastic Optimization) algorithm of Gerber and Douc (2021) for finding the global minimizer of a function defined through an expectation. Informally speaking, G-PFSO can be seen as a particle and derivative-free version of stochastic gradient methods.

**License** GPL (>= 2)

**Imports** Rcpp (>= 1.0.7), Rdpack

**LinkingTo** Rcpp

**RdMacros** Rdpack

**RoxygenNote** 7.1.2

## R topics documented:

PFoptim-package . . . . .	2
gpfs . . . . .	2
SSP_Resampler . . . . .	5
Stratified_Resampler . . . . .	6
<b>Index</b>	<b>8</b>

PFoptim-package

*The 'PFoptim' package: summary information***Description**

The package provides an implementation of the G-PFSO (Global Particle Filter Stochastic Optimization) algorithm of Gerber and Douc (2021) for finding the global minimizer of a function defined through an expectation. In addition, a function for implementing the SSP resampling algorithm (Gerber et al. 2019) and a function for implementing the Stratified resampling algorithm are also provided.

**Author(s)**

Mathieu Gerber

Maintainer: Mathieu Gerber &lt;mathieu.gerber@bristol.ac.uk&gt;

**References**

Gerber M, Chopin N, Whiteley N (2019). "Negative association, ordering and convergence of resampling methods." *The Annals of Statistics*, **47**(4), 2236–2260.

Gerber M, Douc R (2021). "A global stochastic optimization particle filter algorithm." *arXiv preprint arXiv:2007.04803*.

gpfs

*Global Particle filter Stochastic Optimization***Description**

This function implements the G-PFSO (Global Particle Filter Stochastic Optimization) algorithm of Gerber and Douc (2021) for minimizing either the function  $\theta \mapsto E[\text{fn}(\theta, Y)]$  from i.i.d. realizations  $y_1, \dots, y_n$  of  $Y$  or the function  $\theta \mapsto \sum_{i=1}^n \text{fn}(\theta, y_i)$ , where  $\theta$  is a vector of dimension  $d$ .

**Usage**

```
gpfs(obs, N, fn, init, numit, resampling=c("SSP", "STRAT", "MULTI"), ..., control= list())
```

**Arguments**

obs	Either a vector of observations or a matrix of observations (the number of rows being the sample size).
N	Number of particles. The parameter N must be greater or equal to 2.

<code>fn</code>	function for a single observation. If <code>theta</code> is an $N$ by $d$ matrix and <code>y</code> is a single observation (i.e. <code>y</code> is a scalar if <code>obs</code> is a vector and a vector if <code>obs</code> is a matrix) then <code>fn(theta,y)</code> must be a vector of length $N$ . If some rows of <code>theta</code> are outside the search space then the corresponding entries of the vector <code>fn(theta,y)</code> must be equal to <code>Inf</code> .
<code>init</code>	Either a vector of size $d$ or a function used to sample the initial particles such that <code>init(N)</code> is an $N$ by $d$ matrix (or alternatively a vector of length $N$ if $d=1$ ). If <code>init</code> is a vector then the initial distribution is a Gaussian distribution with mean <code>init</code> and covariance matrix equal to <code>sigma_init^2</code> times the identity matrix. By default <code>sigma_init</code> is equal to two. (The value of <code>sigma_init</code> can be changed using the <code>control</code> argument, see below.)
<code>numit</code>	Number of iterations of the algorithm. If <code>numit</code> is not specified then G-PFSO estimates the minimizer of the function $E[\text{fn}(\theta, Y)]$ (in which case the observations are processed sequentially and <code>numit</code> is equal to the sample size). If <code>numit</code> is specified then G-PFSO computes the minimizer of the function $\sum_{i=1}^n \text{fn}(\theta, y_i)$ .
<code>resampling</code>	Resampling algorithm to be used. Resampling should be either "SSP" (SSP resampling), "STRAT" (stratified resampling) or "MULTI" (multinomial resampling).
<code>...</code>	Further arguments to be passed to <code>fn</code> .
<code>control</code>	A list of control parameters. See details.

## Details

Note that arguments after `...` must be matched exactly.

G-PFSO computes two estimators of the minimizer of the objective function, namely the estimators  $\hat{\theta}_{\text{numit}}^N$  and  $\tilde{\theta}_{\text{numit}}^N$ . The former is defined by  $\hat{\theta}_{\text{numit}}^N = \frac{1}{\text{numit}} \sum_{t=1}^{\text{numit}} \hat{\theta}_t^N$  and converges to a particular element of the search space at a faster rate than the latter, but the latter estimator can find more quickly a small neighborhood of the minimizer of the objective function.

By default the sequence  $(t_p)_{p \geq 0}$  is taken as

$$t_p = t_{p-1} + \lceil \max(A t_{p-1}^\varrho \log(t_{p-1}), B) \rceil$$

with  $A=B=1$ ,  $\varrho = 0.1$  and  $t_0 = 5$ . The value of  $A, B, \varrho$  and  $t_0$  can be changed using the `control` argument (see below).

The `control` argument is a list that can supply any of the following components:

**sigma\_init:** Variance parameter of the distribution used by default to sample the initial particles.

**alpha:** Parameter  $\alpha$  of the learning rate  $t^{-\alpha}$ , which must be a strictly positive real number. By default, `alpha=0.5`.

**Sigma:** Scale matrix used to sample the particles. `Sigma` must be either a  $d$  by  $d$  covariance matrix or a strictly positive real number. In this latter case the scale matrix used to sample the particles is `diag(Sigma, d)`. By default, `Sigma=1`.

**trace:** If `trace=TRUE` then the value of  $\tilde{\theta}_t$  and of the effective sample size  $ESS_t$  for all  $t = 1, \dots, \text{numit}$  are returned. By default, `trace=FALSE`.

**indep:** If indep=TRUE and Sigma is a diagonal matrix or a scalar then the Student's t-distributions have independent components. By default, indep=FALSE and if Sigma is a not a diagonal matrix this parameter is ignored.

**A:** Parameter A of the sequence  $(t_p)_{p \geq 0}$  used by default (see above). This parameter must be strictly positive.

**B:** Parameter B of the sequence  $(t_p)_{p \geq 0}$  used by default (see above). This parameter must non-negative.

**varrho:** Parameter varrho of the sequence  $(t_p)_{p \geq 0}$  used by default (see above). This parameter must be in the interval (0,1).

**t0:** Parameter  $t_0$  of the the sequence  $(t_p)_{p \geq 0}$  used by default (see above). This parameter must be a non-negative integer.

**nu:** Number of degrees of freedom of the Student's t-distributions used at time  $t \in (t_p)_{p \geq 0}$  to generate the new particles. By default nu=10

**c\_ess:** A resampling step is performed when  $ESS_t \leq N_{c_{ess}}$ . This parameter must be in the interval (0,1] and by default c\_ess=0.7.

## Value

A list with the following components:

B_par	Value of $\bar{\theta}_{\text{numit}}^N$
T_par	Value of $\hat{\theta}_{\text{numit}}^N$
T_hist	Value of $\hat{\theta}_t^N$ for $t = 1, \dots, \text{numit}$ (only if trace=TRUE)
ESS	Value of the effective sample for $t = 1, \dots, \text{numit}$ (only if trace=TRUE)

## References

Gerber M, Douc R (2021). "A global stochastic optimization particle filter algorithm." *arXiv preprint arXiv:2007.04803*.

## Examples

```
#Definition of fn
fn_toy<-function(theta, obs){
  test<-rep(0,nrow(theta))
  test[theta[,2]>0]<-1
  ll<-rep(-Inf,nrow(theta))
  ll[test==1]<-dnorm(obs,mean=theta[test==1,1], sd=theta[test==1,2],log=TRUE)
  return(-ll)
}
#Generate data y_1,...,y_n
n<-10000 #sample size
theta_star<-c(0,1) #true parameter value
y<-rnorm(n,mean=theta_star[1], sd=theta_star[2])
d<-length(theta_star)
#Define init function to be used
pi0<-function(N){
  return(cbind(rnorm(N,0,5), rexp(N)))
}
```

```

}
##Example 1: Maximum likelihood estimation in the Gaussian model
##true value of the MLE
mle<-c(mean(y),sd(y))
## use gpfso to compute the MLE
Est<-gpfso(y, N=100, fn=fn_toy, init=pi0, numit=20000, control=list(trace=TRUE))
## print \bar{\theta}^{N_{\text{numit}}} and \tilde{\theta}^{N_{\text{numit}}}
print(Est$B_par)
print(Est$T_par)
##assess convergence
par(mfrow=c(1,2))
for(k in 1:2){
  plot(Est$T_hist[,k],type='l', xlab="iteration", ylab="estimated value")
  lines(cumsum(Est$T_hist[,k])/1:length(Est$T_hist[,k]),type='l', col='red')
  abline(h=mle[k])
}
##Example 2: Expected log-likelihood estimation in the Gaussian model
## Estimation of theta_star using gpfso
Est<-gpfso(y, N=100, fn=fn_toy, init=pi0, control=list(trace=TRUE))
## print \bar{\theta}^{N_{\text{numit}}} and \tilde{\theta}^{N_{\text{numit}}}
print(Est$B_par)
print(Est$T_par)
##assess convergence
par(mfrow=c(1,2))
for(k in 1:2){
  plot(Est$T_hist[,k],type='l', xlab="iteration", ylab="estimated value")
  lines(cumsum(Est$T_hist[,k])/1:length(Est$T_hist[,k]),type='l', col='red')
  abline(h=theta_star[k])
}

```

---

SSP\_Resampler

*SSP resampling*


---

## Description

This function implements the SSP resampling algorithm (Gerber et al. 2019).

## Usage

```
SSP_Resampler(U,W)
```

## Arguments

W	A vector of normalized weights.
U	A vector of points in (0,1) such that $\text{length}(U)=\text{length}(W)$ .

## Details

For efficiency reasons, SSP\_Resampler does not perform checks on the supplied arguments.

**Value**

A vector of length  $N$  with elements in the set  $\{1, \dots, N\}$ , with  $N = \text{length}(U) = \text{length}(W)$ .

**References**

Gerber M, Chopin N, Whiteley N (2019). “Negative association, ordering and convergence of resampling methods.” *The Annals of Statistics*, **47**(4), 2236–2260.

**Examples**

```
N<-100
W<-rbeta(N,0.5,2)
W<-W/sum(W)
J<-SSP_Resampler(runif(N),W)
```

---

Stratified_Resampler	<i>Stratified resampling</i>
----------------------	------------------------------

---

**Description**

This function implements the stratified resampling algorithm described see e.g. in Section 9.6 of Chopin and Papaspiliopoulos (2020)

**Usage**

```
Stratified_Resampler(U,W)
```

**Arguments**

W	A vector of normalized weights.
U	A vector of points in (0,1) such that $\text{length}(U) = \text{length}(W)$ .

**Details**

For efficiency reasons, `Stratified_Resampler` does not perform checks on the supplied arguments.

**Value**

A vector of length  $N$  with elements in the set  $\{1, \dots, N\}$ , with  $N = \text{length}(U) = \text{length}(W)$ .

**References**

Chopin N, Papaspiliopoulos O (2020). *An introduction to sequential Monte Carlo*, volume 4. Springer.

**Examples**

```
N<-100  
W<-rbeta(N,0.5,2)  
W<-W/sum(W)  
J<-Stratified_Resampler(runif(N),W)
```

# Index

- \*Topic **algorithms**
  - SSP\_Resampler, [5](#)
  - Stratified\_Resampler, [6](#)
- \*Topic **filters**
  - gpfs, [2](#)
- \*Topic **global**
  - gpfs, [2](#)
- \*Topic **optimization,**
  - gpfs, [2](#)
- \*Topic **particle**
  - gpfs, [2](#)
- \*Topic **resampling**
  - SSP\_Resampler, [5](#)
  - Stratified\_Resampler, [6](#)
- \*Topic **stochastic**
  - gpfs, [2](#)

gpfs, [2](#)

PFoptim (PFoptim-package), [2](#)

PFoptim-package, [2](#)

SSP\_Resampler, [5](#)

Stratified\_Resampler, [6](#)